

Kubernetes - Part 2

6-MONTH INDUSTRY MASTERCLASS PATHWAY

An enterprise-grade, comprehensive technical architecture, deployment workflow, security hardening guide, and observability plan. Designed for senior engineering professionals (6+ years experience) striving for Principal SRE and Cloud Architect roles.

CORE CONCEPTS | PRODUCTION HARDENING | 20+ INCIDENT SCENARIOS

Automated SRE Learn System | Generated Pdf Generator

Kubernetes Deep-Dive Study Guide (Part 2 of 3)

Advanced Configurations, Performance Tuning, Security Capabilities, Sandboxing, and Scale Boundaries

1. Part Introduction and Scope

This study guide focuses on the internal mechanics of Kubernetes. It is designed for engineers seeking to transition from operational competency to architectural mastery.

PART 2 SCOPE				
Advanced Scheduling	Kernel & Kubelet	Advanced Security &		
& Placement	Performance Tuning	Container Sandboxing		
- Topology Spread	- cgroups v2 & Systemd	- eBPF Network Policies		
- NUMA-Aware Scheduling	- CPU/Topology Managers	- gVisor (runsc) Sentry		
- Priority & Preemption	- Eviction Thresholds	- Kata MicroVMs (QEMU)		
- Custom Schedulers	- Sysctl Kernel Tuning	- Seccomp & AppArmor		

Scope of Coverage

- **Advanced Scheduling & Placement Mechanics:** Deep-dive into scheduling cycles, topology-spread constraints, node/pod affinity, taints/tolerations, priority classes, and preemption.
- **Performance Tuning (OS, Kernel, & Kubelet):** Optimizing the OS kernel, configuring `cgroups v2`, tuning `kubelet` via the `KubeletConfiguration` API, managing CPU/Device/Topology managers, and mitigating CFS quota throttling.
- **Advanced Security & Hardening:** Designing eBPF-based zero-trust network policies, implementing Pod Security Standards (PSS), configuring custom `seccomp` profiles, and managing Linux capabilities.
- **Container Sandboxing (gVisor & Kata Containers):** Architecture, integration, and performance trade-offs of user-space kernels (`runsc`) and microVM runtimes (`kata-runtime`).
- **Scale Boundaries & Control Plane Optimization:** Identifying and resolving control plane bottlenecks across the API-server, `etcd`, and the scheduler at scale.

2. Why These Concepts are Critical for High-Availability Systems

In a production environment, standard configurations lead to degradation under load. Understanding advanced configurations and security boundaries is critical for several key reasons:

Prevention of "Noisy Neighbor" Effects & Pod Eviction Cascades

Without explicit CPU pinning (``static`` CPU Manager policy) and proper Quality of Service (QoS) planning, high-throughput workloads experience latency spikes due to CFS (Completely Fair Scheduler) quota throttling and CPU migrations across NUMA sockets.

Improperly configured eviction thresholds can trigger cascading node failures: as one node runs out of memory (OOM) and evicts pods, the remaining nodes absorb the load, exceed their resource limits, and crash sequentially.

Mitigation of Kernel-Level Security Compromises

Standard containers share the host's Linux kernel via the standard container runtime (``runc``). If an attacker exploits a zero-day vulnerability in the host kernel (e.g., Dirty COW, Privilege Escalation bugs), they can break out of the container and gain root access to the underlying node.

Implementing sandboxed runtimes like gVisor (intercepting syscalls in user space) or Kata Containers (isolating workloads inside dedicated hardware-assisted microVMs) creates a strict security boundary for untrusted or multi-tenant code execution.

High-Throughput, Low-Latency Networking

Standard ``kube-proxy`` implementations rely on ``iptables``, which suffers from $O(N)$ lookup latency (where $O(N)$ scales linearly with the number of Services and Endpoints).

In large-scale clusters, this results in significant packet processing overhead and CPU consumption. Transitioning to eBPF-based networking (e.g., Cilium) changes this lookup complexity to $O(1)$, bypassing the TCP/IP stack in the host kernel to route packets directly to container interfaces.

3. Real-World Enterprise Use Cases

Use Case 1: Ultra-Low Latency Financial Order Execution Engine

An algorithmic trading platform requires sub-millisecond end-to-end latency for its execution engine pods. The workloads must run on bare-metal Kubernetes nodes to bypass hypervisor overhead.



```

? [L1/L2 Cache] [L3 Cache] ? ? [L1/L2 Cache] [L3 Cache] ?
? ?????????? ?????????? ? ? ?????????? ?????????? ?
? ? Core0 ? ? Core1 ? ? ? ? Core2 ? ? Core3 ? ?
????????????????????????????????????????????????????????????
? ? ? ?
High-Speed NIC ??[PCIe]??? ?????????????????????? ?? PCIe (High
Latency)
(Direct NUMA 0 Connection)

```

- **Architectural Requirements:**
- **NUMA Node Alignment:** The pods must be scheduled on the same NUMA node where the high-speed network interface card (NIC) is physically connected via PCIe to avoid cross-socket UPI (Ultra Path Interconnect) latency.
- **Dedicated CPU Allocations:** Workloads require exclusive access to physical CPU cores, bypassing the CFS scheduler completely to eliminate context-switching overhead.
- **Sandboxing:** Unvetted third-party quantitative analysis scripts must run on the same infrastructure but must be sandboxed to prevent intellectual property theft and host kernel access.
- **Solution Architecture:**
- Kubelet configured with `cpuManagerPolicy: static` and `topologyManagerPolicy: single-numa-node`.
- Pods declared as **Guaranteed QoS** (matching CPU/Memory requests and limits) with integer CPU requests.
- Integration of **gVisor** (`runsc`) as a `RuntimeClass` for the untrusted scripting workloads, while the core trading engine runs on native `runc` with CPU pinning.

Use Case 2: Multi-Tenant SaaS Platform with Hard Isolation (PCI-DSS & HIPAA Compliant)

A multi-tenant healthcare enterprise hosts multiple independent clients on a single shared EKS cluster.

```

????????????????????????????????????????????????????????????????
? Multi-Tenant EKS ?
????????????????????????????????????????????????????????????????
?
????????????????????????????????????????????????????????????????
? ?
????????????????????????????????????????????????????????????????
? Tenant A Namespace ? Tenant B Namespace ?
????????????????????????????????????????????????????????????????
? - RuntimeClass: gVisor ? ? - RuntimeClass: Kata ?
? - NetworkPolicy: Deny [Blocked Layer 3] ? - NetworkPolicy: Deny ?
? - Seccomp: Custom Whitel? ? - Seccomp: Custom Whitel?
????????????????????????????????????????????????????????????????

```

- **Architectural Requirements:**
- **Data Isolation:** Zero cross-tenant network or storage access.
- **Kernel Segregation:** A vulnerability in Tenant A's application must not allow access to Tenant B's data


or the underlying host node.

- **Strict Scheduling:** Tenant workloads must be distributed evenly across Availability Zones (AZs) to prevent single-AZ failures from causing tenant-wide outages, while ensuring no single node hosts more than two pods from the same tenant.
- **Solution Architecture:**
- **Network Isolation:** eBPF-based network policies (Cilium) enforcing default-deny ingress and egress rules, with identity-aware L7 filtering.
- **Compute Isolation:** Pods are assigned to specific node groups using `nodeAffinity` and isolated from other tenants using `topologySpreadConstraints` with `maxSkew: 1` across `topology.kubernetes.io/zone`.
- **Runtime Security:** Tenant pods are forced to run inside **Kata Containers** (using AWS bare-metal `metal` instances or Nitro Enclaves that support nested virtualization) via a mutating webhook that injects the appropriate `RuntimeClass`.

4. Comprehensive Architecture Explanation

The interaction between the Kubernetes Control Plane, the Kubelet, the Container Runtime Interface (CRI), and low-level virtualization/sandboxing engines is structured as a layered execution pipeline.

Component Interaction Flow



?	(Shared Ring0) ?	????????????????
?	?	? Gofer (FS I/O) ?
????????????????		????????????????

The Scheduling Cycle

When a Pod is submitted, the `kube-scheduler` executes two primary phases: Scheduling (Filtering and Scoring) and Binding.

1. Filtering (Predicates): Evaluates node resources, ports, taints, tolerations, and affinity rules.
2. Scoring (Priorities): Ranks the remaining nodes based on topology spread, image locality, and node affinity weightings.
3. Reserve & Permit: Reserves node resources in-memory before writing to the API Server.
4. Binding: Writes a `Binding` object to the API Server, which updates the Pod's `spec.nodeName`.

The Kubelet Performance Pipeline

Upon detecting a Pod assigned to its node, the `kubelet` coordinates resource allocation via its internal subsystems:

- **Cgroup Manager:** Configures the Linux `cgroups v2` hierarchy under `/sys/fs/cgroup/kubepods.slice/` to enforce CPU, Memory, and I/O limits.
- **CPU Manager:** If configured with the `static` policy, it allocates exclusive physical CPUs to pods in the **Guaranteed** QoS class, moving their PIDs out of the shared pool.
- **Topology Manager:** Consults the CPU Manager and Device Manager (e.g., GPU allocations) to align allocations to a single NUMA node, avoiding cross-socket interconnect bottlenecks.

Container Runtime Interface (CRI) & Sandbox Execution

The `kubelet` sends a gRPC request to `containerd` (or `CRI-O`) to create the pod sandbox and containers.

- **Standard Runtime (runc):** `containerd` spawns a `containerd-shim-runc-v2` process, which calls `runc` to configure namespaces (namespaces, mount, uts, ipc, pid, user, cgroup) and launch the container process directly on the host kernel.
- **Sandboxed Runtime (gVisor - runsc):** `containerd` spawns `containerd-shim-runsc-v2`. It launches the **Sentry** (a user-space kernel written in Go) and the **Gofer** (a file system proxy). The application container runs inside this sandbox; all system calls are intercepted and handled by the Sentry, preventing direct execution on the host kernel.
- **Sandboxed Runtime (Kata Containers):** `containerd` spawns `containerd-shim-kata-v2`, which launches a lightweight virtual machine (using QEMU, Cloud Hypervisor, or Firecracker). Inside this microVM, a minimal Linux kernel boots, running its own nested container runtime to execute the workload.

5. Types, Classifications, and Components

1. Advanced Scheduling Primitives

- **Topology Spread Constraints:** Controls how pods are distributed across failure domains (zones, regions, nodes) to achieve high availability and resource balance.
- **Node Affinity/Anti-Affinity:** Directs pods to specific nodes based on labels, using hard (`requiredDuringSchedulingIgnoredDuringExecution`) or soft (`preferredDuringSchedulingIgnoredDuringExecution`) constraints.
- **Pod Affinity/Anti-Affinity:** Co-locates or segregates pods relative to other pods on the same node or topology domain.
- **Taints & Tolerations:** Allows nodes to repel certain pods, ensuring only authorized workloads are scheduled on specialized hardware (e.g., GPU nodes, control planes).

2. Container Runtime Classifications

- **Native OCI Runtimes (`runc`):** Direct host kernel execution. Highest performance, lowest security isolation boundary.
- **User-Space Kernel Emulators (gVisor - `runsc`):** Intercepts syscalls. Moderate performance overhead (due to syscall translation), high security isolation boundary.
- **MicroVM Hypervisors (Kata, Firecracker):** Hardware-assisted virtualization. High isolation, higher memory overhead, slower boot times compared to native runtimes.

CONTAINER RUNTIME TAXONOMY			
Runtime	Isolation Model	Performance vs.	Security
runc	Namespaces/cgroups	[Perf: ???]	[Security: ???]
gVisor (runsc)	Syscall Emulation	[Perf: ???]	[Security: ???]
Kata Containers	MicroVM (QEMU/FC)	[Perf: ???]	[Security: ???]

3. Linux Control Group Versions (cgroups)

- **cgroups v1:** Resource controllers are managed independently, leading to synchronization issues between memory and block I/O controllers.
- **cgroups v2:** Single unified hierarchy. Allows accurate resource accounting, better OOM-kill handling, and rootless container execution.

6. Step-by-Step Production Implementation Guide

This guide details the process of configuring a production-grade worker node with a dual-runtime configuration (`runc` + `gVisor`) and configuring the `kubelet` for low-latency performance tuning.

Worker Node Configuration Flow

```

????????????????????????????????????????????????????????????????????????????????????
? 1. Install gVisor  ? 2. Configure  ? 3. Tune Kernel  ?
?   Binaries         ?   ?   containerd  ?   ?   (sysctl.conf)  ?
????????????????????????????????????????????????????????????????????????????????????
????????????????????????????????????????????????????????????????????????????????????
?
????????????????????????????????????????????????????????????????????????????????????
? 6. Apply Runtime  ? 5. Configure  ?
?   Class in K8s    ?   ?   Kubelet    ?
????????????????????????????????????????????????????????????????????????????????????

```

Step 1: Install gVisor (`runsc`) on the Kubernetes Worker Node

Execute the following commands on your worker nodes to install `runsc` and its container-shim:

```

# 1. Download the gVisor binaries
(
  set -e
  ARCH=$(uname -m)
  URL="https://storage.googleapis.com/gvisor/releases/release/latest/${ARCH}"
  wget ${URL}/runsc ${URL}/runsc.sha256 ${URL}/containerd-shim-runsc-v2
  ${URL}/containerd-shim-runsc-v2.sha256

  # Verify checksums
  sha256sum -c runsc.sha256
  sha256sum -c containerd-shim-runsc-v2.sha256

  # Install binaries to system path
  chmod a+rx runsc containerd-shim-runsc-v2
  sudo mv runsc containerd-shim-runsc-v2 /usr/local/bin/
)

```

Step 2: Configure `containerd` to Register the `runsc` Runtime

Modify `/etc/containerd/config.toml` to register the `runsc` handler. Ensure you place this under the correct CRI runtime configuration path:

```

version = 2

[plugins]
[plugins."io.containerd.grpc.v1.cri"]
  [plugins."io.containerd.grpc.v1.cri".containerd]
    default_runtime_name = "runc"

```



```
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes]
  # Standard Native Runtime
  [plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc]
    runtime_type = "io.containerd.runc.v2"
    [plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc.options]
      SystemdCgroup = true

  # Sandboxed gVisor Runtime
  [plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runsc]
    runtime_type = "io.containerd.runsc.v2"
```

Restart `containerd` to apply the configuration:

```
sudo systemctl restart containerd
```

Step 3: Optimize Host Kernel Parameters (`/etc/sysctl.conf`)

Add the following high-performance, low-latency kernel configurations to the host OS:

```
# Prevent packet drops under high connection spikes
net.core.somaxconn = 32768
net.ipv4.tcp_max_syn_backlog = 16384

# Increase the maximum number of open files (file descriptors)
fs.file-max = 2097152

# Adjust virtual memory configuration for heavy database workloads
vm.max_map_count = 262144
vm.overcommit_memory = 1

# Disable Swap to ensure consistent Kubelet scheduling
vm.swappiness = 0

# Optimize ARP cache table sizes for large-scale clusters
net.ipv4.neigh.default.gc_thresh1 = 2048
net.ipv4.neigh.default.gc_thresh2 = 4096
net.ipv4.neigh.default.gc_thresh3 = 8192
```

Apply the sysctl parameters:

```
sudo sysctl --system
```

Step 4: Configure the Kubelet for Performance Tuning

Create or update `/var/lib/kubelet/config.yaml` to enable the `static` CPU Manager policy, set Topology Manager constraints, and configure eviction thresholds:

```
apiVersion: kubelet.config.k8s.io/v1beta1
```

```
kind: KubeletConfiguration
cpuManagerPolicy: "static"
cpuManagerReconcilePeriod: "10s"
topologyManagerPolicy: "single-numa-node"
memoryManagerPolicy: "Static"
evictionHard:
  memory.available: "500Mi"
  nodefs.available: "10%"
  nodefs.inodesFree: "5%"
  imagefs.available: "15%"
evictionPressureTransitionPeriod: "5m0s"
systemReserved:
  cpu: "1000m"
  memory: "1Gi"
  ephemeral-storage: "2Gi"
kubeReserved:
  cpu: "1000m"
  memory: "1Gi"
  ephemeral-storage: "2Gi"
enforceNodeAllocatable:
  - "pods"
```

Restart the `kubelet`:

```
sudo systemctl restart kubelet
```

7. Standard CLI Commands with Deep Technical Explanations

1. Inspecting the Kubelet CPU Manager State

```
kubectrl exec -it <worker-node-ssh-or-pod> -- cat /var/lib/kubelet/cpu_manager_state
```

- **Explanation:** This command outputs the internal state file of the Kubelet CPU Manager. It shows the active CPU assignments, verifying if the `static` policy has successfully isolated cores for **Guaranteed** pods and created the default pool for other workloads.

2. Checking Node NUMA Layout and CPU Topology

```
lscpu | grep -E "NUMA|Thread|Core|Socket"
```

- **Explanation:** Displays the hardware layout of the host, including threads per core, cores per socket, and NUMA nodes. This information is critical for designing Topology Spread Constraints and validating NUMA-aware scheduling.

3. Querying Container Runtime via `crictl`

```
sudo crictl -r unix:///run/containerd/containerd.sock pods
sudo crictl -r unix:///run/containerd/containerd.sock inspect <container-id> | grep -i runtime
```

- **Explanation:**
- `-r unix:///run/containerd/containerd.sock`: Directs `crictl` to use the local containerd gRPC socket.
- `inspect`: Retrieves the low-level JSON configuration of the runtime container, allowing you to verify if a pod is executing via `runsc` (gVisor) or `runc` (native).

4. Analyzing Pod Scheduling Decisions and Topology Spread Constraints

```
kubectrl get pods -o
custom-columns=NAME:.metadata.name,NODE:.spec.nodeName,STATUS:.status.phase --all-namespaces
```

- **Explanation:** Outputs a clean list of pods and their assigned nodes across all namespaces. This allows you to verify if your scheduling policies (such as Topology Spread Constraints) are distributing pods as expected.

8. Production Configuration Examples

1. Hardened, Performance-Tuned Pod Manifest (`production-workload.yaml`)

This manifest configures a high-performance, sandboxed container with CPU pinning, zone-based distribution, and kernel-level security hardening.

```
apiVersion: v1
kind: Pod
metadata:
  name: dynamic-api-gateway
  namespace: core-ingress
  labels:
    app.kubernetes.io/name: api-gateway
    app.kubernetes.io/part-of: core-platform
spec:
  # Enforce execution in the gVisor sandbox
  runtimeClassName: gvisor

  # Priority class to protect against preemption
  priorityClassName: high-priority-preemption

  # Anti-affinity and Topology Spread to guarantee high-availability
  affinity:
    podAntiAffinity:
```

```
requiredDuringSchedulingIgnoredDuringExecution:
  - labelSelector:
      matchExpressions:
        - key: app.kubernetes.io/name
          operator: In
          values:
            - api-gateway
      topologyKey: "kubernetes.io/hostname"

topologySpreadConstraints:
  - maxSkew: 1
    topologyKey: "topology.kubernetes.io/zone"
    whenUnsatisfiable: DoNotSchedule
    labelSelector:
      matchExpressions:
        - key: app.kubernetes.io/name
          operator: In
          values:
            - api-gateway

containers:
  - name: gateway-engine
    image: envoyproxy/envoy:v1.28.0
    imagePullPolicy: IfNotPresent

# Strict Resource Requests and Limits (Guaranteed QoS)
resources:
  limits:
    cpu: "4"          # Integer CPU triggers static CPU pinning
    memory: "8Gi"
  requests:
    cpu: "4"
    memory: "8Gi"

# Hardened Security Context
securityContext:
  allowPrivilegeEscalation: false
  readOnlyRootFilesystem: true
  runAsNonRoot: true
  runAsUser: 10001
  runAsGroup: 10001
  capabilities:
    drop:
      - ALL
    add:
      - NET_BIND_SERVICE # Allow binding to privileged ports (<1024)
  seccompProfile:
    type: RuntimeDefault
```

```

ports:
  - containerPort: 443
    name: https
    protocol: TCP

volumeMounts:
  - name: ephemeral-tmp
    mountPath: /tmp

volumes:
  - name: ephemeral-tmp
    emptyDir:
      medium: Memory # Store temp files in RAM to avoid disk I/O bottlenecks

```

2. RuntimeClass Definition ('gvisor-runtimeclass.yaml')

Registers the gVisor runtime with the Kubernetes cluster control plane.

```

apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: gvisor
handler: runsc # Maps directly to the containerd config.toml handler registration
scheduling:
  nodeSelector:
    sandbox-enabled: "true" # Ensure sandboxed pods only land on prepared nodes

```

3. PriorityClass Definition ('high-priority-preemption.yaml')

Configures a priority tier to ensure critical workloads can preempt lower-priority pods during resource constraints.

```

apiVersion: scheduling.k8s.io/v1
kind: PriorityClass
metadata:
  name: high-priority-preemption
value: 1000000
globalDefault: false
description: "Mission-critical API Gateway services that require immediate scheduling."
preemptionPolicy: PreemptLowerPriority

```

9. Security Considerations & Hardening Best Practices

????????????????????????????

```

? Pod Security Standards ?
????????????????????????????
? - Privileged (None)    ?
? - Baseline (Minimal)   ?
? - Restricted (Strict)  ?
????????????????????????
?
????????????????????????????????????????????????????????????
?
????????????????????????????????????????????????????????????
? Host Hardening        ?
????????????????????????????????????????????????????????????
? - Read-Only Root FS   ?
? - Drop Capabilities    ?
? - Seccomp Profiles     ?
????????????????????????????????????????????????????????????
? Network Isolation     ?
????????????????????????????????????????????????????????????
? - Default-Deny eBPF   ?
? - Namespace Separation ?
? - DNS Query Filtering  ?
????????????????????????????????????????????????????????????

```

1. Pod Security Standards (PSS) Compliance

Kubernetes provides built-in isolation tiers. Production clusters should enforce the Restricted profile at the namespace level:

```

apiVersion: v1
kind: Namespace
metadata:
  name: secure-workloads
  labels:
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/enforce-version: latest

```

The Restricted profile enforces the following constraints:

- Prevents privilege escalation (`allowPrivilegeEscalation: false`).
- Requires containers to run as a non-root user.
- Restricts access to the host network, PID, and IPC namespaces.
- Disallows volume mounts from the host filesystem (`hostPath`).

2. Linux Capabilities Hardening

By default, the container runtime grants a subset of Linux kernel capabilities to containers (e.g., `CAP\_NET\_RAW`, `CAP\_MKNOD`). Attackers can leverage these capabilities to execute container breakouts.

Always drop all default capabilities and selectively add only those explicitly required by the application:

```

securityContext:
  capabilities:
    drop:

```

```
- ALL
add:
- NET_BIND_SERVICE
```

3. eBPF-Based Network Isolation (Cilium NetworkPolicy)

Standard `NetworkPolicies` lack visibility into Layer 7 protocols. Using eBPF-based systems like Cilium allows you to enforce identity-aware policies that filter traffic up to the application layer (L7), including DNS-level filtering:

```
apiVersion: "cilium.io/v2"
kind: CiliumNetworkPolicy
metadata:
  name: secure-api-egress
  namespace: secure-workloads
spec:
  endpointSelector:
    matchLabels:
      app: api-gateway
  egress:
    # Allow egress only to specific DNS endpoints
    - toFQDNs:
        - matchName: "api.stripe.com"
      toPorts:
        - ports:
            - port: "443"
              protocol: TCP
    # Allow DNS queries only to Kube-DNS
    - toEndpoints:
        - matchLabels:
            "k8s:k8s-app": kube-dns
      toPorts:
        - ports:
            - port: "53"
              protocol: ANY
    rules:
      dns:
        - matchPattern: "*"
---
```

10. Observability & Monitoring Considerations

Key Prometheus Metrics to Monitor

Metric Name	Type	Target Threshold	Technical Significance
-------------	------	------------------	------------------------

| :--- | :--- | :--- | :--- |

| `container\_cpu\_cfs\_throttled\_seconds\_total` | Counter | $\$ > 0$ indicates throttling | Measures the total time a container's execution was throttled by the kernel due to CFS quotas. |

| `kube\_pod\_container\_status\_waiting\_reason` | Gauge | Alert on `OOMKilled` | Tracks pods stuck in waiting states. Indicates memory limits are set too low relative to peak usage. |

| `etcd\_disk\_wal\_fsync\_duration\_seconds` | Histogram | $\$ < 10 \text{ms}$ (99th percentile) | Measures latency of writing transaction logs to disk. Values $\$ > 10 \text{ms}$ can lead to etcd leader election timeouts. |

| `apiserver\_request\_duration\_seconds` | Histogram | $\$ < 1 \text{s}$ for non-list calls | Monitors control plane API latency. Spikes indicate resource exhaustion on the control plane nodes. |

Log Aggregation for Sandboxed Runtimes

Monitoring sandboxed containers requires specialized log ingestion configurations:

- **gVisor (`runsc`):** The `Sentry` generates its own internal logs regarding intercepted system calls. These logs should be redirected to a dedicated audit logging pipeline using the `runsc` flag `--debug-log` to capture potential intrusion attempts.
- **Kata Containers:** Standard standard output (`stdout`) and standard error (`stderr`) streams are routed through the `kata-runtime` shim. Since Kata runs inside a virtual machine, kernel logs (`dmesg`) from inside the microVM must be aggregated via the `kata-agent` and forwarded to the host's log collector (e.g., Fluentbit, Vector) to diagnose driver or kernel panics inside the sandbox.

11. Common Troubleshooting Scenarios with RCA (Root Cause Analysis)

Scenario 1: Severe CPU Throttling Despite Low Average CPU Usage

- **Symptom:** A microservice's average CPU usage sits at 20%, but its response latency is highly degraded. The metric `container\_cpu\_cfs\_throttled\_seconds\_total` is rapidly incrementing.
- **Root Cause Analysis (RCA):** This is caused by the multi-threaded nature of the application and the CFS quota enforcement mechanism. The Linux kernel evaluates CPU quotas over a period (typically 100ms). If an application spawns 10 threads that execute simultaneously, it can consume its entire 100ms quota within the first 10ms of the period. The kernel then throttles the container for the remaining 90ms.
- **Resolution Steps:**

1. Increase the CPU limit to accommodate multi-threaded bursts, or set CPU requests and limits to the same integer value (Guaranteed QoS) to trigger CPU pinning.
2. If running on supported kernels, disable CFS quota limits entirely via the Kubelet configuration (`cpuCFSQuota: false`) while retaining resource requests for scheduling.

Scenario 2: Pod Stuck in `Terminating` State Due to Storage Unmount Failures

- **Symptom:** A pod is deleted but remains stuck in the `Terminating` state indefinitely.
- **Root Cause Analysis (RCA):** The Kubelet is unable to unmount the volume associated with the pod because the mount point is marked as busy. This often occurs when a container process is blocked in an uninterruptible sleep state (D-state) due to a network partition with the underlying storage controller (e.g., NFS, AWS EBS, or Ceph). Because the process cannot be terminated by `SIGKILL`, the container runtime cannot release the mount.
- **Resolution Steps:**

1. Identify the blocked process on the worker node:

```
ps aux | grep 'D '
```

2. Inspect the mount path and verify the storage controller's network connectivity:

```
df -h | grep <pod-uid>
```

3. Force-unmount the blocked filesystem on the host node:

```
umount -f -l  
/var/lib/kubelet/pods/<pod-uid>/volumes/kubernetes.io~csi/<volume-name>/mount
```

Scenario 3: `TopologyAffinityError` During Scheduling

- **Symptom:** A pod is rejected during scheduling with the error `TopologyAffinityError: Resources cannot be allocated due to Topology Affinity constraints`.
- **Root Cause Analysis (RCA):** The Kubelet's Topology Manager is configured with the `single-numa-node` policy. The pod requested a combination of resources (e.g., 8 CPUs and 1 GPU) that cannot be physically allocated from a single NUMA node. For example, the worker node may have 4 free CPUs on NUMA Node 0 and 4 free CPUs on NUMA Node 1, or the GPU may be physically aligned to NUMA Node 0 while all free CPUs reside on NUMA Node 1.
- **Resolution Steps:**

1. Check the node's resource distribution:

```
numactl --hardware
```

2. Adjust the pod's resource requests so they fit within the physical boundaries of a single NUMA node.

3. Change the Kubelet's Topology Manager policy to `best-effort` if strict alignment is not mandatory for the workload.

12. Common Mistakes and How to Avoid Them in Production

Mistake 1: Relying on Memory Limits Without Accounting for Kernel Page Cache

Setting memory limits too close to the application's heap size can cause sudden container terminations.

```

????????????????????????????????????????????????????????????????????????????????????
?                                MEMORY ALLOCATION OVERHEAD                                ?
????????????????????????????????????????????????????????????????????????????????????
? JVM Heap / Application Memory          ? Kernel Page Cache / I/O Buffers?
? (Explicitly controlled by application)? (Grows dynamically during I/O) ?
????????????????????????????????????????????????????????????????????????????????????
? ????????????????????????????????????????? Total Pod Memory ????????????????????????? ?
????????????????????????????????????????????????????????????????????????????????????

```

- **The Mistake:** Configuring a JVM application with `-Xmx4g`` and setting the Kubernetes memory limit to `4Gi``.
- **Why it fails:** The JVM process requires memory beyond the heap, including thread stacks, metaspace, and native memory allocations. Additionally, the kernel allocates memory for page caches during file I/O operations. If the combined memory usage exceeds `4Gi``, the kernel's OOM-killer terminates the container.
- **The Solution:** Ensure the Kubernetes memory limit is at least 20-30% higher than the application's heap limit. For Java applications, utilize the dynamic container awareness options:

```
-XX:MaxRAMPercentage=75.0
```

Mistake 2: Over-committing CPU on Nodes Running Latency-Sensitive Workloads

- **The Mistake:** Over-provisioning CPU resources by scheduling too many pods with low CPU requests and high limits on the same node.
- **Why it fails:** When the node experiences a traffic spike, all containers attempt to burst up to their limits. The CPU scheduler is forced to context-switch between hundreds of active threads, causing scheduling latency to spike from microseconds to hundreds of milliseconds.
- **The Solution:** Group latency-sensitive workloads into dedicated node pools using node taints. Set their QoS class to **Guaranteed** (requests equal to limits) to pin them to dedicated physical cores.

13. Enterprise-Level Recommendations

1. Networking: eBPF Routing with Cilium

At scale, enterprise clusters should replace `kube-proxy`` with Cilium in eBPF-replacement mode. This bypasses the host's `iptables`` rules entirely, routing packets directly from the container's network namespace to the physical interface. This optimization reduces network latency by up to 30% and frees up CPU cycles on the host nodes.

2. Etcd Optimization for High-Density Clusters

To support high-frequency scheduling and scaling operations, `etcd`` must be optimized:

- **Disk Isolation:** Mount the `etcd`` write-ahead log (WAL) on a dedicated, physical NVMe drive to prevent

write latency spikes caused by shared disk I/O.

- ****Compaction Tuning:**** Configure the API Server to perform auto-compaction of the resource history every 5 minutes:

```
--etcd-compaction-interval=5m
```

14. Advanced Concepts

1. Kubelet Eviction Policies (Soft vs. Hard)

The Kubelet acts as the last line of defense against node instability. It monitors host resources and can evict pods to reclaim memory, disk space, or inodes.

Resource Depletion			
		?	
????????????????????		????????????????????	
?		?	
????????????????????		????????????????????	
? Hard Threshold ?		? Soft Threshold ?	
? (Immediate Evict) ?		? (Grace Period) ?	
????????????????????		????????????????????	
? - No grace period ?		? - Wait for grace ?	
? - SIGKILL sent ?		? - SIGTERM first ?	
? - Fast recovery ?		? - Graceful exit ?	
????????????????????		????????????????????	

- ****Hard Eviction (`evictionHard`):**** If a resource drops below the threshold, the Kubelet immediately terminates the target pods without a grace period. It sends a `SIGKILL` signal to the container processes to reclaim resources as quickly as possible.
- ****Soft Eviction (`evictionSoft`):**** If a resource drops below the threshold, the Kubelet waits for a user-configured grace period (`evictionSoftGracePeriod`) before initiating eviction. Pods are terminated gracefully using `SIGTERM` first.

2. cgroups v2 Architecture & Memory Management Improvements

The transition to `cgroups v2` introduces two key memory control files that improve workload stability:

- ``memory.max``: The hard limit. Exceeding this value triggers the host OOM-killer immediately.
- ``memory.high``: A soft limit/throttle threshold. When a container exceeds this value, the kernel slows down the container's allocation rate by forcing it into direct page reclamation. This allows the container to continue running while releasing unused memory, reducing sudden OOM-kills.

15. Integration with Other DevOps Tools

```

Git Repository
?
? (ArgoCD)
????????????????????
?   Kubernetes Cluster ?
????????????????????
?
????????????????????????????????????????
?
????????????????????????????????????????
? Node Provisioning ?           ? Node OS Tuning ?
? (Terraform / ASG) ?           ?   (Ansible)   ?
????????????????????????????????????????
? - Custom UserData ?           ? - Install runsc ?
? - Configure Kubelet ?           ? - Apply sysctl ?
????????????????????????????????????????

```

1. Terraform (Infrastructure as Code)

When provisioning managed Kubernetes nodes (e.g., AWS EKS), configure the `kubelet` performance flags using custom launch templates and user data:

```

resource "aws_launch_template" "eks_tuned_nodes" {
  name_prefix   = "eks-tuned-"
  image_id      = data.aws_ami.eks_optimized.id
  instance_type = "c6i.4xlarge" # Compute-optimized instance with dual NUMA nodes

  user_data = base64encode(<<-EOF
    #!/bin/bash
    set -o xtrace
    # Modify KubeletConfiguration prior to bootstrap
    mkdir -p /etc/kubernetes/kubelet/
    cat <<KUBE_CFG > /etc/kubernetes/kubelet/kubelet-config.json
    {
      "cpuManagerPolicy": "static",
      "topologyManagerPolicy": "single-numa-node",
      "systemReserved": {"cpu": "1000m", "memory": "1Gi"},
      "kubeReserved": {"cpu": "1000m", "memory": "1Gi"}
    }
    KUBE_CFG

    # Bootstrap EKS Node
    /etc/eks/bootstrap.sh my-cluster-name \
      --kubelet-extra-args '--config=/etc/kubernetes/kubelet/kubelet-config.json'
  EOF
)

```

```
)
}
```

2. Ansible (Configuration Management)

Use Ansible to automate gVisor installation and kernel sysctl tuning across bare-metal or self-managed nodes:

```
- name: Configure High-Performance Sandboxed Worker Nodes
  hosts: kubernetes_workers
  become: true
  tasks:
    - name: Download gVisor runsc binary
      get_url:
        url: https://storage.googleapis.com/gvisor/releases/release/latest/x86_64/runsc
        dest: /usr/local/bin/runsc
        mode: '0755'

    - name: Configure sysctl performance parameters
      sysctl:
        name: "{{ item.key }}"
        value: "{{ item.value }}"
        state: present
        reload: yes
      with_items:
        - { key: 'net.core.somaxconn', value: '32768' }
        - { key: 'vm.max_map_count', value: '262144' }
```

16. Comparison Tables with Competing Tools

The following table compares container runtimes across key operational metrics:

Feature / Metric	Native OCI (`runc`)	gVisor (`runsc`)	Kata Containers	WebAssembly (Wasmtime)
Isolation Mechanism	Namespaces & Cgroups	User-space Syscall Emulation	MicroVM (Hypervisor)	WebAssembly Sandbox Engine
Syscall Overhead	Extremely Low (< 1%)	Moderate to High (15-30%)	Low to Moderate (5-10%)	Minimal (< 2%)
Memory Footprint	Minimal (< 15MB)	Low (15-30MB)	High (100MB-256MB per VM)	Extremely Low (< 5MB)
Boot Time	< 100ms	~150ms	1-2s	< 10ms
Kernel Version	Shared with Host	Emulated (independent)	Dedicated guest kernel	None (WASI abstraction)

|

| Best Use Case | Trusted Internal Services | Untrusted SaaS code execution | Hard multi-tenancy, legacy kernels | High-performance serverless, edge |

17. Visual Cheat Sheet

Advanced Scheduling & Runtime Operations Reference

```
????????????????????????????????????????????????????????????????????????????????????
?                                SCHEDULING & RUNTIME CHEAT SHEET                                ?
????????????????????????????????????????????????????????????????????????????????????
?  KUBELET CONFIGURATION (/var/lib/kubelet/config.yaml)                                     ?
?  ??? cpuManagerPolicy: "static"                -> Enables exclusive CPU pinning.        ?
?  ??? topologyManagerPolicy: "single-numa"      -> Aligns CPU/PCIe/Device to NUMA.        ?
?  ??? evictionHard: "memory.available<500Mi"   -> Prevents node kernel panic.          ?
????????????????????????????????????????????????????????????????????????????????????
?  CONTAINERD RUNTIMES (/etc/containerd/config.toml)                                     ?
?  ??? [runtimes.runc] -> Native execution, shared host kernel.                        ?
?  ??? [runtimes.runsc] -> Sandboxed execution, intercepts syscalls (gVisor).          ?
????????????????????????????????????????????????????????????????????????????????????
?  CLI DIAGNOSTICS                                                                    ?
?  ??? crictl inspect <id> | grep -i runtime -> Check active container type.            ?
?  ??? cat /var/lib/kubelet/cpu_manager_state -> Verify CPU core assignments.            ?
?  ??? numactl --hardware                        -> Inspect host NUMA topology.          ?
????????????????????????????????????????????????????????????????????????????????????
?  POD QUALITY OF SERVICE (QoS)                                                        ?
?  ??? Guaranteed -> requests == limits (Trigger for static CPU pinning)                ?
?  ??? Burstable  -> requests != limits (Eligible for resource overcommit)              ?
?  ??? BestEffort -> No requests/limits (First candidate for eviction)                  ?
????????????????????????????????????????????????????????????????????????????????????
```

18. Comprehensive Final Learning Summary

To master advanced Kubernetes configurations, prioritize these key takeaways:

1. Scheduling Control: Move beyond simple node selectors. Utilize Topology Spread Constraints to distribute workloads across zones, and combine them with PriorityClasses to ensure critical pods can preempt lower-priority workloads during resource constraints.
2. Deterministic Performance: For high-throughput or low-latency workloads, bypass default scheduling behavior.

Configure the Kubelet with a ``static`` CPU Manager policy and a ``single-numa-node`` Topology Manager policy. Ensure your pods are in the Guaranteed QoS class (matching integer CPU requests and limits) to pin them to dedicated physical cores and align them with local PCIe devices.

3. Defense in Depth: Enforce the Restricted Pod Security Standard across all production namespaces. For workloads running untrusted code or operating in multi-tenant environments, implement container sandboxing. Use gVisor (``runsc``) to intercept and handle system calls in user space, or Kata Containers to isolate workloads inside dedicated microVMs.

4. Platform Observability: Set up proactive alerts for key system metrics. Monitor ``container_cpu_cfs_throttled_seconds_total`` to identify CPU throttling issues, and track ``etcd_disk_wal_fsync_duration_seconds`` to catch storage latency bottlenecks before they impact control plane stability.

Q21. Pod Sandboxing (gVisor vs. Kata Containers vs. Firecracker) via RuntimeClass

Detailed Answer:

By default, Kubernetes uses the host kernel shared across all containers via namespaces and cgroups. In multi-tenant or untrusted code execution environments, this shared-kernel model presents a significant security risk (kernel privilege escalation, container escapes). Pod Sandboxing mitigates this by introducing a secondary isolation layer using a ``RuntimeClass`` resource mapped to specialized Container Runtime Interface (CRI) shims.

There are three primary sandboxing architectures:

1. gVisor (``runsc``): A user-space kernel written in Go. It intercepts system calls made by the application and implements them inside user-space, filtering and translating them before they ever reach the host kernel (via ``ptrace`` or KVM). It has a low memory footprint but introduces significant system call latency overhead (bad for I/O-heavy workloads).
2. Kata Containers (``kata-runtime``): A lightweight Virtual Machine (VM) execution layer. Each Pod runs inside its own dedicated hardware-isolated VM with its own kernel. It uses QEMU, Cloud Hypervisor, or Firecracker as its Virtual Machine Monitor (VMM). It provides strong hardware-level isolation but has a higher memory/CPU footprint and slower startup times than gVisor.
3. Firecracker (via MicroVMs): An open-source VMM written in Rust, specifically designed for serverless workloads. It boots minimal virtual machines in milliseconds. In Kubernetes, it is typically orchestrated via Kata Containers (using the Firecracker hypervisor plugin) or custom CRI implementations like ``firecracker-containerd``.

To configure sandboxing, the cluster administrator must define a ``RuntimeClass`` and configure the underlying CRI daemon (e.g., ``containerd``) to handle the corresponding handler.

Production Scenario / Practical Example:

An SRE team needs to run untrusted, user-submitted Python data-science scripts on a shared Kubernetes cluster. They decide to run these pods inside gVisor (``runsc``) to prevent kernel-level exploits.

First, configure ``/etc/containerd/config.toml`` on the worker nodes:

```
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runsc]
```

```
runtime_type = "io.containerd.runsc.v1"
```

Restart containerd:

```
sudo systemctl restart containerd
```

Next, define the `RuntimeClass` in Kubernetes:

```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: gvisor
handler: runsc
```

Apply the `RuntimeClass` and reference it in the untrusted Pod specification:

```
apiVersion: v1
kind: Pod
metadata:
  name: untrusted-python-runner
  namespace: tenant-sandbox
spec:
  runtimeClassName: gvisor
  containers:
  - name: script-executor
    image: python:3.11-slim
    command: ["python", "-c", "import os; print(os.uname())"]
    resources:
      limits:
        cpu: "1"
        memory: "512Mi"
      requests:
        cpu: "200m"
        memory: "256Mi"
```

Q22. Advanced Scheduling: Custom Schedulers and Scheduler Framework Extension Points

Detailed Answer:

The default Kubernetes Scheduler (`kube-scheduler`) executes a 2-phase cycle to assign a Pod to a Node: the Scheduling Cycle (selects a node) and the Binding Cycle (commits the decision to the API server). The modern `kube-scheduler` is designed around the Scheduler Framework, which exposes extension points as Go APIs.

Developers can compile custom scheduling logic directly into the scheduler binary or run multiple concurrent schedulers inside the cluster. The extension points run in the following sequence:

1. QueueSort: Sorts the scheduling queue of pending Pods.
2. PreFilter: Prepares or checks information about the Pod before filtering.
3. Filter: Excludes nodes that cannot run the Pod (equivalent to Predicates in the old scheduler).
4. PostFilter: Invoked if no schedulable nodes are found (used for preemption/back-off).
5. PreScore: Performs pre-scoring work to generate state for scoring.
6. Score: Ranks nodes for the filtered Pod (equivalent to Priorities in the old scheduler).
7. NormalizeScore: Modifies scores before the scheduler calculates the final node ranking.
8. Reserve: Reserves resources on the selected node to avoid race conditions (optimistic binding).
9. Permit: Approves, denies, or delays the binding (can hold a Pod until dependent pods are ready-ideal for gang scheduling).
10. PreBind: Performs work required before a Pod is bound (e.g., provisioning network volumes).
11. Bind: Binds the Pod to the Node in the API Server.
12. PostBind: Cleans up resources or triggers telemetry after binding succeeds.

To run a custom scheduler alongside the default one, you configure a `KubeSchedulerConfiguration` deployment.

Production Scenario / Practical Example:

An SRE team deploys a custom scheduler called `batch-priority-scheduler` that implements co-scheduling (gang scheduling) for high-performance computing (HPC) workloads using the `Permit` plugin.

Below is the `KubeSchedulerConfiguration` config map used to initialize the custom scheduler:

```
apiVersion: kubescheduler.config.k8s.io/v1
kind: KubeSchedulerConfiguration
leaderElection:
  leaderElect: true
  resourceName: batch-priority-scheduler
  resourceNamespace: kube-system
profiles:
  - schedulerName: batch-priority-scheduler
    plugins:
      permit:
        enabled:
          - name: "CoschedulingPlugin"
```

Deploying a Pod that targets this custom scheduler:

```
apiVersion: v1
kind: Pod
metadata:
  name: hpc-worker-node-0
```

```
labels:
  pod-group: hpc-simulation-run-42
spec:
  schedulerName: batch-priority-scheduler
  containers:
  - name: mpi-worker
    image: mpi-operator:latest
    resources:
      limits:
        cpu: "4"
        memory: "8Gi"
```

Q23. API Server Performance Tuning: Flow Schema and Priority Level Configuration (APF)

Detailed Answer:

Historically, the Kubernetes API Server had limited mechanisms to protect itself from being overwhelmed by API requests (e.g., buggy controller loops, massive scale deployments, or runaway scripts). It relied on a simple max-in-flight limit (`--max-requests-inflight` and `--max-mutating-requests-inflight`). If reached, all subsequent requests were dropped, regardless of priority.

API Priority and Fairness (APF) solves this by categorizing incoming requests into distinct flows and managing them using queues. APF is governed by two Custom Resources:

1. **FlowSchema**: Classifies incoming requests (based on user, group, namespace, verb, or resource) and maps them to a specific `PriorityLevelConfiguration`.
2. **PriorityLevelConfiguration**: Defines the resource allocation (concurrency shares) and queuing behavior for classified requests. It specifies how many concurrent execution slots (concurrency limits) are allocated and how requests are queued (using shuffle sharding to isolate abusive clients) if concurrency limits are exceeded.

The total concurrency capacity of the API Server is divided among priority levels using "Nominal Concurrency Shares." If a priority level does not use its allocated shares, other active levels can borrow them.

Production Scenario / Practical Example:

In a large cluster, a rogue custom operator is hammering the API Server with `GET` and `LIST` requests, starving critical system controllers. The SRE team creates a dedicated `FlowSchema` and `PriorityLevelConfiguration` to throttle this specific service account while guaranteeing resources for system components.

Create the `PriorityLevelConfiguration` for low-priority background operators:

```
apiVersion: flowcontrol.apiserver.k8s.io/v1beta3
kind: PriorityLevelConfiguration
metadata:
  name: background-operator-limit
```

```
spec:
  type: Queue
  queue:
    queues: 64
    handSize: 4
    queueLengthLimit: 50
  limited:
    nominalConcurrencyShares: 10
    lendablePercent: 0
    limitResponse:
      type: Queue
```

Create the `FlowSchema` matching the rogue operator's ServiceAccount (`rogue-operator-sa`) and map it to the newly created priority level:

```
apiVersion: flowcontrol.apiserver.k8s.io/v1beta3
kind: FlowSchema
metadata:
  name: throttle-rogue-operator
spec:
  priorityLevelConfiguration:
    name: background-operator-limit
  matchingPrecedence: 500
  distinguisherMethod:
    type: ByUser
  rules:
  - subjects:
    - kind: ServiceAccount
      serviceAccount:
        name: rogue-operator-sa
        namespace: production
    resourceRules:
    - verbs: ["get", "list", "watch"]
      apiGroups: ["*"]
      resources: ["*"]
```

Q24. Kernel Tuning via Sysctls (Safe vs. Unsafe Sysctls)

Detailed Answer:

Kubernetes allows containerized applications to modify kernel parameters (`sysctls`) at runtime. However, because the kernel is shared, modifying certain parameters can destabilize the host or interfere with other containers on the same node.

Kubernetes divides sysctls into two categories:

1. Safe Sysctls: These are namespaces-isolated parameters. Modifying them inside a container does not impact other containers or the host. Examples include:

- ``kernel.shm_rmid_forced``
- ``net.ipv4.ip_local_port_range``
- ``net.ipv4.tcp_keepalive_time``

2. Unsafe Sysctls: These parameters are not fully isolated or can have global side effects on host stability, networking, or memory management. Examples include:

- ``net.core.somaxconn`` (increases maximum socket backlog)
- ``net.ipv4.tcp_tw_reuse`` (enables fast socket recycling)
- ``kernel.msgmax`` (sets maximum message size in bytes)

To use unsafe sysctls, the cluster administrator must explicitly whitelist them in the ``kubelet`` configuration file on each node. If an unsafe sysctl is requested in a Pod spec but not whitelisted in the kubelet, the Pod will fail to schedule or fail to start with a ``SysctlForbidden`` error.

Production Scenario / Practical Example:

An SRE team is running a high-throughput Nginx ingress controller that is dropping connections under heavy traffic due to socket backlog exhaustion. They must increase ``net.core.somaxconn`` (unsafe) and adjust ``net.ipv4.ip_local_port_range`` (safe).

First, edit the Kubelet configuration file (``/var/lib/kubelet/config.yaml``) on all worker nodes to whitelist the unsafe sysctl:

```
apiVersion: kubelet.config.k8s.io/v1beta1
kind: KubeletConfiguration
allowedUnsafeSysctls:
  - "net.core.somaxconn"
```

Restart the Kubelet daemon:

```
sudo systemctl restart kubelet
```

Now, deploy the Nginx Pod specifying both the safe and unsafe sysctls in the ``securityContext``:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: high-throughput-nginx
  namespace: ingress
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
```

```

metadata:
  labels:
    app: nginx
spec:
  securityContext:
    sysctls:
      - name: net.ipv4.ip_local_port_range
        value: "1024 65535"
      - name: net.core.somaxconn
        value: "32768"
  containers:
    - name: nginx
      image: nginx:1.25-alpine
      resources:
        limits:
          cpu: "2"
          memory: "2Gi"
        requests:
          cpu: "500m"
          memory: "512Mi"

```

Q25. Node Resource Reservation: `system-reserved`, `kube-reserved`, and `eviction-hard`

Detailed Answer:

To prevent system instability, memory starvation, and SSH unresponsiveness under heavy container workloads, Kubernetes SREs must configure Node Resource Reservations. By default, if not configured, the kubelet assumes that all node resources are available for scheduling Pods.

The total capacity of a node is allocated according to the following equation:

$$Allocatable = Capacity - kube-reserved - system-reserved - eviction-threshold$$

Where:

- **`kube-reserved`**: Resources reserved for Kubernetes system daemons (kubelet, container runtime, node-problem-detector, etc.).
- **`system-reserved`**: Resources reserved for host OS system daemons (sshd, udev, systemd, journald, etc.).
- **`eviction-hard`**: Thresholds (such as memory availability or disk space) that, when crossed, trigger the kubelet to evict pods to preserve host stability.

If these reserves are not configured, a memory leak in a pod can cause the OS Out-Of-Memory (OOM) killer to terminate critical system processes (like `containerd` or `kubelet`), taking the entire node offline.

Production Scenario / Practical Example:

An SRE tunes a bare-metal worker node with 128 GB of RAM and 32 CPU cores. They need to reserve resources to ensure the host remains responsive even under maximum workload stress.

Edit the `/var/lib/kubelet/config.yaml` on the node:

```
apiVersion: kubelet.config.k8s.io/v1beta1
kind: KubeletConfiguration
kubeReserved:
  cpu: "1000m"
  memory: "2Gi"
  ephemeral-storage: "5Gi"
systemReserved:
  cpu: "500m"
  memory: "1.5Gi"
  ephemeral-storage: "5Gi"
evictionHard:
  memory.available: "500Mi"
  nodefs.available: "10%"
  nodefs.inodesFree: "5%"
evictionMinimumReclaim:
  memory.available: "100Mi"
```

Verify the node's allocatable capacity after restarting the kubelet:

```
kubectl get node node-01-bm -o jsonpath='{.status.allocatable}'
```

The output will show reduced memory and CPU capacities compared to `.status.capacity`, matching the reservation rules.

Q26. Advanced Network Policies: DNS-Based Policies and Layer 7 Enforcement

Detailed Answer:

Standard Kubernetes `NetworkPolicy` resources operate exclusively at OSI Layer 3 (IP addresses, CIDR blocks) and Layer 4 (TCP, UDP ports). They cannot enforce rules based on domain names (FQDNs) or Layer 7 protocols (HTTP paths, methods, gRPC services).

To solve this, advanced Container Network Interfaces (CNIs) like Cilium leverage eBPF (Extended Berkeley Packet Filter) and Envoy sidecars/proxies to provide Layer 7 visibility and DNS-based egress filtering.

- ****DNS-based policies****: Allow pods to only resolve and connect to specific external domains (e.g., `api.stripe.com`), dynamically updating the underlying IP whitelist as DNS records change.
- ****Layer 7 policies****: Allow or deny traffic based on HTTP fields (e.g., allow `GET /v1/public`, deny `POST /v1/admin`).

Using standard network policies for external APIs is fragile because public cloud services dynamically change their IP addresses, causing static CIDR blocks to break.

Production Scenario / Practical Example:

An SRE needs to secure a payment processing microservice (`payment-processor`). It must *only* be allowed to:

1. Resolve external DNS queries for `api.stripe.com`.
2. Connect to `api.stripe.com` over HTTPS (port 443).
3. Restrict internal ingress traffic to only allow `GET /health` and `POST /charge` endpoints from the `frontend` service.

This is implemented using a `CiliumNetworkPolicy`:

```
apiVersion: "cilium.io/v2"
kind: CiliumNetworkPolicy
metadata:
  name: payment-processor-rules
  namespace: finance
spec:
  endpointSelector:
    matchLabels:
      app: payment-processor
  egress:
    # Rule 1: Allow DNS resolution
    - toEndpoints:
        - matchLabels:
            "k8s:io.kubernetes.pod.namespace": kube-system
            k8s-app: kube-dns
      toPorts:
        - ports:
            - port: "53"
              protocol: UDP
        rules:
          dns:
            - matchPattern: "*"
    # Rule 2: Allow dynamic egress to Stripe FQDN on port 443
    - toFQDNs:
        - matchName: "api.stripe.com"
      toPorts:
        - ports:
            - port: "443"
              protocol: TCP
  ingress:
    # Rule 3: Allow L7 ingress from frontend
    - fromEndpoints:
        - matchLabels:
            app: frontend
      toPorts:
        - ports:
            - port: "8080"
```

```
    protocol: TCP
  rules:
    http:
      - method: "GET"
        path: "/health"
      - method: "POST"
        path: "/charge"
```

Q27. Container Runtime Interface (CRI) Performance Tuning in `containerd`

Detailed Answer:

Under high-concurrency workloads (e.g., high-churn serverless architectures, batch jobs starting thousands of containers simultaneously), the Container Runtime Interface (CRI) can become a bottleneck. In most modern clusters, `containerd` serves as the CRI.

To prevent container creation delays, timeout failures, and runtime API unresponsive errors, several performance configurations in `/etc/containerd/config.toml` must be tuned:

- **max_container_ops**: Limits the number of concurrent container operations. For high-churn nodes, this must be scaled up.
- **registry.mirrors**: Configures local registries or pull-through caches to prevent Docker Hub rate-limiting and minimize pull latency.
- **gRPC Configuration**: Tuning the max message sizes and connection timeouts prevents the kubelet-to-containerd gRPC channel from dropping under load.
- **io.containerd.grpc.v1.cri.containerd.runtimes**: Assigning optimized configurations to the default runc runtime, such as enabling systemd cgroups (`SystemdCgroup = true`), which prevents double-bookkeeping of cgroup hierarchies and improves host resource tracking.

Production Scenario / Practical Example:

An SRE tunes `containerd` on a node group dedicated to processing large-scale batch jobs where up to 200 containers are spun up within 10 seconds.

Optimize the `/etc/containerd/config.toml` file:

```
version = 2

[plugins."io.containerd.grpc.v1.cri"]
  sandbox_image = "registry.k8s.io/pause:3.9"
  max_container_ops = 500
  max_concurrent_downloads = 20

[plugins."io.containerd.grpc.v1.cri".containerd]
  disable_snapshot_annotations = true
```



```
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc]
  runtime_type = "io.containerd.runc.v2"

[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc.options]
  SystemdCgroup = true

[plugins."io.containerd.grpc.v1.cri".registry]
  config_path = "/etc/containerd/certs.d"
[plugins."io.containerd.grpc.v1.cri".registry.mirrors."docker.io"]
  endpoint = ["https://mirror.gcr.io", "https://registry-1.docker.io"]
```

Apply changes:

```
sudo systemctl restart containerd
```

Q28. Securing the Data Plane: Migrating to Pod Security Standards (PSS) and Admission Controllers

Detailed Answer:

Pod Security Policies (PSP) were deprecated in Kubernetes 1.21 and completely removed in 1.25. The modern replacement is Pod Security Standards (PSS), built directly into the Kubernetes API Server, and managed via namespace labels.

PSS defines three distinct profiles:

1. Privileged: Unrestricted policy. Allows root execution, host namespace sharing, and privilege escalation.
2. Baseline: Minimally restrictive policy. Prevents known privilege escalations (disallows host namespaces, host paths, and blocks dangerous capabilities).
3. Restricted: Heavily restrictive policy. Enforces strict hardening (requires running as non-root, dropping all default capabilities except safe ones, disabling write access to root filesystems).

PSS applies these profiles using three actions (modes): `enforce` (blocks non-compliant pods), `audit` (records violations in audit logs), and `warn` (returns a warning message to the client).

For advanced enterprise requirements (e.g., blocking images from untrusted registries, requiring specific labels, enforcing resource limits), SREs deploy Admission Controllers like Kyverno or OPA Gatekeeper.

Production Scenario / Practical Example:

An SRE team needs to secure the `production` namespace. They must enforce the `restricted` Pod Security Standard and deploy a Kyverno policy that blocks containers from running if their image tag is `latest`.

First, label the namespace to enforce the `restricted` PSS:

```
apiVersion: v1
```

```
kind: Namespace
metadata:
  name: production
  labels:
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/enforce-version: "v1.28"
    pod-security.kubernetes.io/warn: restricted
    pod-security.kubernetes.io/warn-version: "v1.28"
```

Next, deploy a Kyverno policy to block the `latest` image tag in the `production` namespace:

```
apiVersion: kyverno.io/v1
kind: ClusterPolicy
metadata:
  name: block-latest-tag
spec:
  validationFailureAction: Enforce
  background: true
  rules:
    - name: validate-image-tag
      match:
        any:
          - resources:
              namespaces:
                - production
              kinds:
                - Pod
      validate:
        message: "Using the 'latest' tag is forbidden in the production namespace."
        pattern:
          spec:
            containers:
              - image: "!*:latest"
```

Q29. Topology-Aware Routing (Topology-Aware Hints)

Detailed Answer:

In multi-AZ Kubernetes clusters, network traffic routing can incur high cross-AZ data transfer costs and add latency (typically 1-2ms per hop across zones). Topology-Aware Routing (formerly Topology-Aware Hints) optimizes this by routing service traffic to endpoints located within the same Availability Zone (AZ) as the originating Pod.

When a Service has Topology-Aware Routing enabled, the EndpointSlice controller generates "hints" on endpoint slices. The local routing component (such as `kube-proxy` or an eBPF-based CNI) inspects these hints and dynamically configures IPVS/iptables/eBPF maps to prioritize endpoints in the same zone.

Conditions for Activation:

To prevent overloading a single zone, Topology-Aware Routing automatically disables itself if endpoints are unevenly distributed across zones. Specifically, it requires:

- The ratio of endpoints in a zone must roughly match the ratio of allocatable CPU resources of the nodes in that zone.
- A minimum number of endpoints (typically ≥ 3 endpoints per zone) to ensure redundancy.

Production Scenario / Practical Example:

An SRE runs a latency-sensitive microservices stack in AWS across three zones (`us-east-1a`, `us-east-1b`, `us-east-1c`). They must configure the `order-service` to route traffic to the `inventory-service` within the same AZ.

Annotate the `inventory-service` to activate Topology-Aware Routing:

```
apiVersion: v1
kind: Service
metadata:
  name: inventory-service
  namespace: commerce
  annotations:
    service.kubernetes.io/topology-mode: Auto
spec:
  selector:
    app: inventory
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```

Verify that the EndpointSlice controller has generated the topology hints:

```
kubectl get endpointslice -l kubernetes.io/service-name=inventory-service -o yaml
```

Look for the `hints` block in the output:

```
endpoints:
  - addresses:
      - 10.244.1.42
    conditions:
      ready: true
    zone: us-east-1a
    hints:
      forZones:
        - name: us-east-1a
```

Q30. Envelope Encryption with KMS Provider v2

Detailed Answer:

By default, Kubernetes Secrets are stored in `etcd` in plaintext (base64 encoded). If etcd backups are leaked or an attacker gains access to the etcd volume, all secrets are compromised. Envelope Encryption secures etcd by encrypting secret data with a local Data Encryption Key (DEK), which is itself encrypted using a Key Encryption Key (KEK) managed by an external Key Management Service (KMS) like AWS KMS, Azure Key Vault, or HashiCorp Vault.

KMS v2 (introduced as stable in v1.29) brings major performance and architectural improvements over KMS v1:

- **Key Rotation**: KMS v2 automates key rotation without requiring an API server restart.
- **Performance**: KMS v2 uses a metadata-driven caching mechanism that drastically reduces the number of gRPC calls to the external KMS provider, preventing API server throttling during high-scale deployments.
- **Status API**: Introduces health check endpoints to detect KMS availability issues early.

The API Server communicates with a local KMS plugin (running as a daemon on the control plane) via a UNIX domain socket using gRPC.

Production Scenario / Practical Example:

An SRE configures AWS KMS v2 envelope encryption for a production Kubernetes cluster.

First, create the `EncryptionConfiguration` file on the Control Plane node at `/etc/kubernetes/encryption/config.yaml`:

```
apiVersion: apiserver.config.k8s.io/v1
kind: EncryptionConfiguration
resources:
  - resources:
    - secrets
  providers:
    - kms:
      name: aws-kms-provider
      apiVersion: v2
      endpoint: unix:///var/run/kms-provider.sock
      timeout: 3s
    - identity: {} # Fallback to plaintext if KMS fails (not recommended for production
write, but good for migration)
```

Next, update the `/etc/kubernetes/manifests/kube-apiserver.yaml` manifest to mount the UNIX socket and pass the encryption provider flag:

```
spec:
  containers:
    - name: kube-apiserver
      command:
        - kube-apiserver
```

```
- --encryption-provider-config=/etc/kubernetes/encryption/config.yaml
# ... other flags ...
volumeMounts:
- name: encryption-config
  mountPath: /etc/kubernetes/encryption
  readOnly: true
- name: kms-sock
  mountPath: /var/run/kms-provider.sock
volumes:
- name: encryption-config
  hostPath:
    path: /etc/kubernetes/encryption
- name: kms-sock
  hostPath:
    path: /var/run/kms-provider.sock
    type: Socket
```

Once applied, encrypt existing secrets in the cluster to apply the new KMS policy:

```
kubectl get secrets --all-namespaces -o json | kubectl replace -f -
```

Q31. Large-Scale Cluster Boundaries: Scaling to 5,000 Nodes and etcd Tuning

Detailed Answer:

Scaling a Kubernetes cluster to the official scale limit (5,000 nodes, 150,000 total pods) introduces significant performance degradation across etcd, the API Server, and the control plane.

To support this scale, etcd must be aggressively tuned:

1. Disk I/O Latency: etcd is highly sensitive to write latency. It must run on dedicated NVMe SSDs. You should use `ionice` to prioritize etcd disk I/O:

```
ionice -c2 -n0 -p $(pgrep etcd)
```

2. Database Size Limit: Increase the default database quota from 2GB to the maximum safe limit (8GB) to prevent the database from locking up with a ``database space exceeded`` error.

3. Heartbeat and Election Timeouts: In high-scale or multi-region networks, default timeouts (100ms heartbeat, 1000ms election) can cause false leader elections. Increase these values to stabilize the cluster.

4. Defragmentation: Schedule regular defragmentation jobs to reclaim free space from deleted resources.

Additionally, the control plane needs optimizations:

- Scale up `--max-requests-inflight` on the API Server.
- Tune the Controller Manager concurrent sync loops (e.g., `--concurrent-deployment-syncs=50`).

Production Scenario / Practical Example:

An SRE configures a highly-available, 3-node external etcd cluster to support a 4,000-node cluster.

The etcd systemd unit service file (`/etc/systemd/system/etcd.service`) is configured with optimized parameters:

```
[Unit]
Description=etcd
Documentation=https://github.com/etcd-io/etcd

[Service]
Type=notify
ExecStart=/usr/local/bin/etcd \
  --name=etcd-node-01 \
  --data-dir=/var/lib/etcd \
  --quota-backend-bytes=8589934592 \
  --heartbeat-interval=250 \
  --election-timeout=1250 \
  --auto-compaction-retention=1 \
  --snapshot-count=10000 \
  --listen-peer-urls=https://10.0.1.10:2380 \
  --listen-client-urls=https://10.0.1.10:2379,https://127.0.0.1:2379 \
  --initial-advertise-peer-urls=https://10.0.1.10:2380 \
  --advertise-client-urls=https://10.0.1.10:2379 \

--initial-cluster=etcd-node-01=https://10.0.1.10:2380,etcd-node-02=https://10.0.1.11:2380,etcd-node-03=https://10.0.1.12:2380 \
  --initial-cluster-token=etcd-prod-token \
  --initial-cluster-state=new
Restart=always
RestartSec=5
LimitNOFILE=65536

[Install]
WantedBy=multi-user.target
```

To clear a fragmentation-induced database lockup, the SRE runs:

```
# Defragment the database
ETCDCTL_API=3 etcdctl --endpoints=https://127.0.0.1:2379 \
  --cacert=/etc/etcd/ca.crt --cert=/etc/etcd/server.crt --key=/etc/etcd/server.key \
  defrag

# Disarm the alarm
ETCDCTL_API=3 etcdctl --endpoints=https://127.0.0.1:2379 \
  --cacert=/etc/etcd/ca.crt --cert=/etc/etcd/server.crt --key=/etc/etcd/server.key \
  alarm disarm
```

Q32. Custom Resource Definition (CRD) Performance, Structural Schemas, and Controller Concurrency

Detailed Answer:

Custom Resource Definitions (CRDs) extend Kubernetes but can degrade API Server performance if designed poorly. Modern CRDs require structural schemas (OpenAPI v3 validation schemas). These allow the API Server to perform client-side and server-side validation, pruning of unknown fields, and efficient serialization without loading heavy Go structs.

Key performance optimizations for CRDs and their controllers:

1. Subresources (`/status` and `/scale`): Enabling the status subresource isolates updates to the spec from updates to the status. This avoids triggering unnecessary reconciliation loops when a controller updates a resource's status, preventing infinite execution loops.
2. Conversion Webhooks: If a CRD supports multiple versions (e.g., `v1alpha1` to `v1`), the API Server calls a webhook to convert objects on the fly. This webhook must be highly optimized and fail-fast to prevent API latency.
3. Controller Concurrency: By default, many controllers process events sequentially. Tuning the controller-runtime's `MaxConcurrentReconciles` allows parallel processing of events.
4. Rate-Limiting Workqueues: Utilizing token-bucket or exponential back-off rate limiters prevents the controller from hammering external APIs or the API Server when reconciliation fails.

Production Scenario / Practical Example:

An SRE designs a high-performance CRD for managing tenant databases (`TenantDatabase`) and optimizes the controller's concurrency.

The CRD manifest with `/status` subresource and OpenAPI validation:

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: tenantdatabases.database.infra.io
spec:
  group: database.infra.io
  names:
    kind: TenantDatabase
    listKind: TenantDatabaseList
    plural: tenantdatabases
    singular: tenantdatabase
  scope: Namespaced
  versions:
    - name: v1
      served: true
      storage: true
      subresources:
        status: {}
```

```

schema:
  openAPIV3Schema:
    type: object
    properties:
      spec:
        type: object
        properties:
          engine:
            type: string
            enum: ["postgres", "mysql"]
          sizeGb:
            type: integer
            minimum: 10
            maximum: 1000
          required: ["engine", "sizeGb"]
      status:
        type: object
        properties:
          phase:
            type: string
          connectionString:
            type: string

```

In the controller implementation (Go using `controller-runtime`), the SRE configures the controller options to scale concurrency:

```

import (
    ctrl "sigs.k8s.io/controller-runtime"
    "sigs.k8s.io/controller-runtime/pkg/controller"
)

func (r *TenantDatabaseReconciler) SetupWithManager(mgr ctrl.Manager) error {
    return ctrl.NewControllerManagedBy(mgr).
        For(&databaseinfov1.TenantDatabase{}).
        WithOptions(controller.Options{
            MaxConcurrentReconciles: 20, // Scale concurrent workers
        }).
        Complete(r)
}

```

Q33. Ephemeral Containers for Advanced Production Debugging

Detailed Answer:

In modern cloud-native architectures, security best practices dictate that production container images should be minimal (e.g., distroless or alpine) and run without root privileges. These images do not contain diagnostic tools

like ``curl``, ``nslookup``, ``tcpdump``, or even a shell (``sh`/`bash``).

This makes live debugging difficult when a Pod is failing. Ephemeral Containers solve this. An ephemeral container is a temporary container that you run inside an existing Pod to inspect its state, troubleshoot network configurations, or read local files. It shares the same namespaces (network, pid, ipc) and volumes as the target container but runs with its own diagnostic image.

Because they are temporary, you cannot add them to an existing deployment spec directly (which would trigger a rolling restart). Instead, they are injected dynamically using the ``/ephemeralcontainers`` API subresource.

Production Scenario / Practical Example:

A critical payment microservice running on a distroless image in production is failing to connect to an external database. The SRE needs to debug the network path from inside that specific pod without restarting it.

Run an ephemeral container using a rich Ubuntu-based diagnostic image, sharing the network namespace of the target pod:

```
kubectl debug -it payment-service-7f6984d7-abc12 \
  --image=nicolaka/netshoot \
  --target=payment-container
```

Once inside the interactive shell (provided by the ``netshoot`` image), the SRE can run diagnostics directly within the target container's context:

```
# Check DNS resolution
nslookup db-prod.internal.domain

# Trace the TCP path
termshark -i eth0

# Inspect local processes of the payment-container (since --target was specified, PID
namespace is shared)
ps aux
```

Q34. Graceful Node Shutdown and Priority-Based Shutdown

Detailed Answer:

When a Kubernetes worker node terminates abruptly (e.g., during cloud provider spot-instance preemption, autoscaling down, or physical hardware maintenance), running pods are terminated abruptly. This can cause data corruption, dropped connections, and API errors.

Graceful Node Shutdown allows the ``kubelet`` to detect a system shutdown event (via systemd logind inhibiting locks) and delay the shutdown. This delay gives the kubelet time to cleanly terminate pods using the standard pod lifecycle (sending ``SIGTERM``, running ``preStop`` hooks, and waiting for the termination grace period).

This is configured using two parameters:

- ``shutdownGracePeriod``: The total time the node delays shutdown to allow all pods to terminate.
- ``shutdownGracePeriodCriticalPods``: The portion of the shutdown grace period reserved for critical pods (pods with ``system-node-critical`` or ``system-cluster-critical`` priority classes), ensuring system-level daemons terminate last.

Production Scenario / Practical Example:

An SRE configures a cluster using AWS Spot Instances (which give a 2-minute preemption warning). They must configure the kubelet to gracefully shut down within 90 seconds, reserving 30 seconds of that time specifically for critical system pods.

Edit ``/var/lib/kubelet/config.yaml`` on all spot nodes:

```
apiVersion: kubelet.config.k8s.io/v1beta1
kind: KubeletConfiguration
shutdownGracePeriod: "90s"
shutdownGracePeriodCriticalPods: "30s"
```

To ensure systemd communicates the shutdown event to the kubelet, ensure the ``systemd-logind`` service is running, and verify the kubelet log output during a test shutdown:

```
sudo systemctl status systemd-logind
```

When a shutdown is initiated, the kubelet logs will show:

```
"Graceful node shutdown in progress"
"Evicting pods before node shutdown"
```

Non-critical pods are given 60 seconds (90s - 30s) to gracefully terminate, after which critical pods are terminated during the remaining 30 seconds.

Q35. Service Mesh mTLS vs. CNI-Layer Encryption (WireGuard/IPsec)

Detailed Answer:

Securing data-in-transit within a Kubernetes cluster is a common compliance requirement (e.g., PCI-DSS, HIPAA). There are two primary architectural approaches to encrypting pod-to-pod communication:

Feature	Service Mesh mTLS (e.g., Istio, Linkerd)	CNI-Layer Encryption (e.g., Cilium WireGuard/IPsec)
Layer of Operation	Layer 7 (Application/Session Layer)	Layer 3/4 (Network/Transport Layer)
Implementation	Sidecar proxies (Envoy) or ambient node-proxies intercepting TCP connections.	Kernel-level tunnels encrypting IP packets directly at the network interface.

| Identity & Auth | SPIFFE/SPIRE cryptographic identities, rich mutual TLS (mTLS) handshake. | Node-to-node or pod-to-pod shared keys or IPsec security associations. |

| Performance Overhead| Higher CPU/Memory footprint (due to user-space proxy context-switching). | Extremely low overhead (WireGuard runs in kernel space, highly optimized). |

| Visibility & Control | High (provides L7 metrics, path-based routing, tracing). | Low (only encrypts the traffic, no application-level metrics). |

SRE Recommendation: Use CNI-layer encryption (WireGuard) if the sole requirement is "encryption at rest and in transit" for compliance, as it has minimal performance impact. Use a Service Mesh if you also require advanced traffic management, distributed tracing, and fine-grained L7 authorization.

Production Scenario / Practical Example:

An SRE team running Cilium as their CNI decides to enable WireGuard encryption cluster-wide to satisfy a compliance audit with zero application code or deployment changes.

To enable WireGuard encryption in Cilium, patch the Cilium configuration:

```
kubectl patch daemonset cilium \
  -n kube-system \
  --type='json' \
  -p='[{"op": "add", "path": "/spec/template/spec/containers/0/args/-", "value":
"--enable-wireguard=true"}]'
```

Verify that WireGuard is active and encrypting traffic between nodes:

```
# Log into a worker node and check WireGuard tunnel status
sudo wg show
```

The output will display the `cilium\_wg0` interface, active peers (other nodes), and transfer bytes, confirming that all pod-to-pod traffic crossing node boundaries is encrypted at the kernel level.

Q36. CPU Management Policies: Static CPU Pinning and NUMA-Aware Scheduling

Detailed Answer:

By default, the Kubernetes CPU Manager uses the `none` policy. The kubelet uses CFS (Completely Fair Scheduler) bandwidth control to throttle containers based on their CPU limits. Under this policy, container processes can be dynamically scheduled across any available CPU core. This causes frequent context switching, L1/L2 cache misses, and NUMA node traversal latency, which degrades performance for latency-sensitive workloads (e.g., databases, telco VNFs, high-frequency trading apps).

To solve this, the CPU Manager can be configured with the `static` policy. When a Pod has:

1. A Quality of Service (QoS) class of Guaranteed (CPU requests equal CPU limits).

2. An integer CPU request (e.g., `2`, `4`, but not `2.5` or `500m`).

The `static` policy allocates exclusive, dedicated CPU cores (pinning via `cpuset`) to that container.

To optimize this further, the Topology Manager coordinates resource assignments across CPU sockets and memory banks (NUMA nodes), ensuring that the pinned CPUs and the container's allocated memory reside on the same physical NUMA socket.

Production Scenario / Practical Example:

An SRE configures a worker node group hosting high-performance Redis instances to use CPU pinning and NUMA alignment.

Configure `/var/lib/kubelet/config.yaml` on the target nodes:

```
apiVersion: kubelet.config.k8s.io/v1beta1
kind: KubeletConfiguration
cpuManagerPolicy: "static"
cpuManagerReconcilePeriod: "10s"
topologyManagerPolicy: "single-numa-node"
reservedSystemCPUs: "0,1" # Reserve cores 0 and 1 for host/kubelet processes
```

Restart the kubelet:

```
sudo systemctl restart kubelet
```

Deploy the Redis Pod as a Guaranteed QoS workload with integer CPU requirements:

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-cache-pinned
  namespace: cache
spec:
  containers:
  - name: redis
    image: redis:7.2-alpine
    resources:
      limits:
        cpu: "2"          # Integer value, matches request -> Guaranteed QoS
        memory: "4Gi"     # Matches request
      requests:
        cpu: "2"
        memory: "4Gi"
```

Q37. Memory Management: Cgroup v2, PSI, and Memory QoS

Detailed Answer:

Kubernetes v1.25 introduced stable support for cgroup v2, which improves resource isolation and memory management over cgroup v1. In cgroup v1, memory limits were enforced strictly using the OOM killer. When a container exceeded its memory limit, it was immediately terminated, even if the host had plenty of free memory.

Cgroup v2 introduces Memory QoS (Quality of Service) using the following parameters:

- ``memory.min``: Hard-guaranteed memory. Memory below this threshold is never reclaimed by the kernel. Maps to ``requests.memory``.
- ``memory.low``: Best-effort guaranteed memory. The kernel will only reclaim memory below this threshold if there is extreme memory pressure.
- ``memory.high``: Throttling threshold. If crossed, the kernel throttles the container's processes and performs aggressive page reclaim before invoking the OOM killer. Maps to ``limits.memory`` (proportional calculation).
- ``memory.max``: Absolute hard limit. Crossing this triggers the OOM killer.

Additionally, cgroup v2 supports Pressure Stall Information (PSI), which provides real-time metrics on how much CPU, memory, and I/O starvation is delaying processes. The kubelet uses PSI to make better eviction decisions before the host locks up.

Production Scenario / Practical Example:

An SRE enables Memory QoS to prevent a memory-heavy Java application from causing immediate OOM kills, allowing the kernel to throttle the application's page allocation first.

Ensure the worker nodes are booting with cgroup v2 enabled (verify via ``mount | grep cgroup``).

Enable the ``MemoryQoS`` feature gate in ``/var/lib/kubelet/config.yaml``:

```
apiVersion: kubelet.config.k8s.io/v1beta1
kind: KubeletConfiguration
featureGates:
  MemoryQoS: true
```

Deploy the Java application:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: java-app
  namespace: production
spec:
  replicas: 2
  selector:
    matchLabels:
      app: java-app
  template:
    metadata:
      labels:
```

```
    app: java-app
  spec:
    containers:
      - name: jvm-worker
        image: openjdk:17-jdk-slim
        resources:
          requests:
            memory: "1Gi" # Maps to cgroup v2 memory.min / memory.low
            cpu: "1"
          limits:
            memory: "2Gi" # Maps to cgroup v2 memory.high / memory.max
            cpu: "2"
```

Q38. eBPF-Based Networking: Bypassing kube-proxy with Cilium

Detailed Answer:

The standard Kubernetes networking model relies on kube-proxy, which uses host-level iptables or IPVS (IP Virtual Server) rules to route Service clusterIPs to backend Pod IPs.

- **The iptables Bottleneck**: iptables rules are evaluated sequentially. In a cluster with thousands of services, the iptables chain grows to tens of thousands of rules. Every network packet must traverse this list sequentially, causing significant CPU utilization and network latency ($O(N)$ lookup complexity).
- **The IPVS Bottleneck**: While IPVS improves lookup performance to $O(1)$ using hash tables, it still relies on the Linux kernel's network stack and connection tracking (`conntrack`), which introduces state table overhead.

eBPF-based networking (pioneered by Cilium) completely bypasses iptables, IPVS, and `conntrack`. eBPF allows running sandboxed programs directly inside the Linux kernel network path (at the driver level via XDP or at the traffic control `tc` layer).

When a packet is sent, Cilium's eBPF programs intercept it and perform a direct lookup in a highly optimized in-kernel BPF map to find the destination pod's IP. It then rewrites the packet header and routes it directly to the target interface, bypassing the entire host network stack. This reduces latency, saves host CPU cycles, and supports massive routing scales with $O(1)$ performance.

Production Scenario / Practical Example:

An SRE team deploys a massive microservices cluster and disables `kube-proxy` entirely, replacing it with Cilium's eBPF-based host-routing.

First, when provisioning the cluster (e.g., via `kubeadm`), skip the installation of `kube-proxy`:

```
apiVersion: kubeadm.k8s.io/v1beta3
kind: ClusterConfiguration
kubeProxy:
  config:
```

```
mode: "ipvs" # Or omit kube-proxy completely during deployment
```

Install Cilium via Helm with `kube-proxy` replacement enabled:

```
helm install cilium cilium/cilium \
  --namespace kube-system \
  --set kubeProxyReplacement=true \
  --set k8sServiceHost=api-server-internal.domain \
  --set k8sServicePort=6443
```

Verify that Cilium has successfully taken over service routing and that no iptables rules are being generated for services:

```
# Exec into a Cilium agent pod and check status
cilium status --compact
```

The output will confirm:

```
`KubeProxyReplacement: True`
```

Q39. Admission Webhook Failure Modes and Cascading Outages

Detailed Answer:

Mutating and Validating Admission Webhooks are powerful tools for enforcing security and injecting sidecars (e.g., service meshes). However, they introduce a critical circular dependency: the API Server must query the webhook pod (which runs inside the cluster) to validate or mutate resources.

If the webhook deployment goes down, becomes slow, or experiences network disruption, the API Server can lock up. Depending on the configuration, this can prevent any new pods (including the webhook pods themselves) from starting, leading to a cascading cluster outage.

To prevent this, SREs must configure webhooks defensively:

1. ``failurePolicy``: Set to ``Ignore`` for non-critical webhooks. If set to ``Fail``, the API Server will reject resource creation if the webhook is unreachable.
2. ``timeoutSeconds``: Keep this value low (e.g., ``1s`` to ``3s``). The default is ``10s``, which can quickly exhaust the API Server's max-in-flight request pool under heavy load.
3. ``namespaceSelector`` / ``objectSelector``: Always exclude system namespaces (like ``kube-system``) to ensure critical control plane components can start even if the webhook is failing.
4. ``reinvocationPolicy``: Controls whether a mutating webhook is called again if subsequent webhooks modify the object.

Production Scenario / Practical Example:

An SRE configures a validating webhook for a policy engine (e.g., OPA Gatekeeper) ensuring it is fail-safe and

does not block emergency cluster recovery operations.

The production-grade `ValidatingWebhookConfiguration`:

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: gatekeeper-validating-webhook-configuration
webhooks:
- name: validation.gatekeeper.sh
  rules:
    - apiGroups: ["*"]
      apiVersions: ["*"]
      operations: ["CREATE", "UPDATE"]
      resources: ["*"]
  clientConfig:
    service:
      name: gatekeeper-webhook-service
      namespace: gatekeeper-system
      path: "/v1/admit"
  # Defensive Tuning 1: Fall back to Ignore if the webhook is down
  failurePolicy: Ignore
  # Defensive Tuning 2: Fail fast to prevent API Server exhaustion
  timeoutSeconds: 3
  # Defensive Tuning 3: Exclude system-critical namespaces from validation
  namespaceSelector:
    matchExpressions:
      - key: kubernetes.io/metadata.name
        operator: NotIn
        values: ["kube-system", "gatekeeper-system", "kube-node-lease"]
  sideEffects: None
  admissionReviewVersions: ["v1"]
```

Q40. IPv4/IPv6 Dual-Stack Networking Configuration and Routing Architecture

Detailed Answer:

As enterprise networks run out of private IPv4 addresses (RFC 1918 ranges), especially when scaling large Kubernetes deployments across hybrid clouds, IPv4/IPv6 Dual-Stack networking becomes necessary. Dual-stack allows the allocation of both IPv4 and IPv6 addresses to Pods and Services.

Architectural Requirements:

1. Host Support: Worker nodes must have both IPv4 and IPv6 addresses configured on their primary network interfaces.
2. CNI Support: The CNI (e.g., Calium, Calico) must support dual-stack IPAM (IP Address Management) to

allocate both address families to Pod network interfaces.

3. Kube-Controller-Manager: Must be configured with separate CIDR blocks for IPv4 and IPv6 pods and services.

Service IP Family Policies:

When creating Services in a dual-stack cluster, you configure the `ipFamilyPolicy` field:

- `SingleStack`: Allocates a single IP (family matches the first configured block).
- `PreferDualStack`: Allocates both IPv4 and IPv6 IPs if the cluster is dual-stack; otherwise, falls back to single-stack.
- `RequireDualStack`: Allocates both IPv4 and IPv6 IPs. If dual-stack is not available, the Service creation fails.

CoreDNS Impact: CoreDNS must be configured to return both `A` (IPv4) and `AAAA` (IPv6) records to allow pods to resolve services over either protocol.

Production Scenario / Practical Example:

An SRE configures a dual-stack cluster using `kubeadm` and deploys a dual-stack web service.

The `ClusterConfiguration` manifest used to initialize the control plane:

```
apiVersion: kubeadm.k8s.io/v1beta3
kind: ClusterConfiguration
networking:
  podSubnet: "10.244.0.0/16,2001:db8:42:1::/64"
  serviceSubnet: "10.96.0.0/16,2001:db8:42:2::/112"
```

Initialize the cluster:

```
kubeadm init --config=kubeadm-config.yaml
```

Deploy a dual-stack Nginx Service that requires both IPv4 and IPv6 endpoints:

```
apiVersion: v1
kind: Service
metadata:
  name: dual-stack-web
  namespace: production
spec:
  selector:
    app: web-server
  ipFamilyPolicy: RequireDualStack
  ipFamilies:
    - IPv4
    - IPv6
  ports:
    - protocol: TCP
      port: 80
```

```
    targetPort: 80
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web-server
  namespace: production
spec:
  replicas: 2
  selector:
    matchLabels:
      app: web-server
  template:
    metadata:
      labels:
        app: web-server
    spec:
      containers:
        - name: nginx
          image: nginx:1.25-alpine
```

Verify the service has been allocated two IPs (one IPv4 and one IPv6):

```
kubectl get svc dual-stack-web -n production
```

The output will display:

```
`CLUSTER-IP: 10.96.120.45, 2001:db8:42:2::bc12`
```